

# DharmeshGajera week8 assignment

## Question 1.

### What is UBER mode?

**Ans:** - Normally mappers and reducers will run by Resource Manager (RM), RM will create separate container for mapper and reducer. Uber configuration, will allow to run mapper and reducers in the same process as the Application Master (AM).

If you have a small dataset or you want to run MapReduce on small amount of data, Uber configuration will help you out, by reducing additional time that MapReduce normally spends in mapper and reducers phase.


Uber jobs are jobs that are executed within the MapReduce Application Master. Rather than communicate with RM to create the mapper and reducer containers. The AM runs the map and reduce tasks within its own process and avoided the overhead of launching and communicate with remote containers.

## Question 2.

### Multi-node cluster

**Ans:** -

1. In About section it shows when Resource Manager Started and other details like Resource Manager version and Hadoop version.



## About the Cluster

Logged in as: dr.who

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics	
Apps Submitted	60103
Apps Pending	0
Apps Running	14
Apps Completed	60089
Containers Running	50
Memory Used	86 GB
Memory Total	151.00 GB
Memory Reserved	0 B
VCores Used	50
VCores Total	90
VCores Reserved	0

Cluster Nodes Metrics	
Active Nodes	3
Decommissioning Nodes	0
Decommissioned Nodes	0
Lost Nodes	1
Unhealthy Nodes	0
Rebooted Nodes	0
Shutdown Nodes	0

Scheduler Metrics	
Scheduler Type	Capacity Scheduler
Scheduling Resource Type	[memory-mb (unit=Mi), vcores]
Minimum Allocation	<memory:1024, vCores:1>
Maximum Allocation	<memory:8192, vCores:4>
Maximum Cluster Application Priority	0

Cluster overview

Cluster ID:

1675999795986

ResourceManager state:

STARTED

ResourceManager HA state:

active

ResourceManager HA zookeeper connection state:

Could not find leader elector. Verify both HA and automatic failover are enabled.

ResourceManager RMStateStore:

org.apache.hadoop.yarn.server.resourcemanager.recovery.NullRMStateStore

ResourceManager started on:

Thu Feb 09 22:29:55 -0500 2023

ResourceManager version:

3.3.0 from aa96f1871bfd858f9bac59cf2a81ec470da649af by brahma source checksum e0a276649f889c15d0e8f08eccd0c10 on 2020-07-06T18:58Z

Hadoop version:

3.3.0 from aa96f1871bfd858f9bac59cf2a81ec470da649af by brahma source checksum 5dc29b802d6ccd77b262ef9d04d19c4 on 2020-07-06T18:44Z


2. In the Nodes tab we can see the all nodes of the cluster with details of Containers, Total Used Memory, Total available Memory, Total Used VCoers, Total Available VCoers.

Rack	Node State	Node Address	Node HTTP Address	Last health-update	Health-report	Containers	Allocation Tags	Mem Used	Mem Avail	VCores Used	VCores Avail	Version
/default-rack	RUNNING	w01.itversity.com:35127	w01.itversity.com:8042	Wed Jun 14 09:09:41 -0400 2023		16		29 GB	21.33 GB	16	14	3.3.0
/default-rack	RUNNING	w03.itversity.com:41791	w03.itversity.com:8042	Wed Jun 14 09:08:55 -0400 2023		17		26 GB	24.33 GB	17	13	3.3.0
/default-rack	RUNNING	w02.itversity.com:46669	w02.itversity.com:8042	Wed Jun 14 09:10:29 -0400 2023		12		20 GB	30.33 GB	12	18	3.3.0

to 3 of 3 entries

FirstPrevious1NextLast

3. In the Application there is RUNNING tab which shows only RUNNING applications.



## RUNNING Applications

Cluster

About  
Nodes  
Node Labels  
Applications  
NEW  
NEW SAVING  
SUBMITTED  
ACCEPTED  
RUNNING  
FINISHED  
FAILED  
KILLED  
Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used
60105	0	16	60089	48	80 GB

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes
3	0	0	1


Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>

Show 20 entries

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	Final Status
<a href="#">application_1675999795986_60674</a>	itv006164	Practice	SPARK		default	0	Wed Jun 14 18:41:10 +0550 2023	Wed Jun 14 18:41:10 +0550 2023	N/A	RUNNING	UNDEFINED
<a href="#">application_1675999795986_60673</a>	itv005355	pyspark-shell	SPARK		default	0	Wed Jun 14 18:39:25 +0550 2023	Wed Jun 14 18:39:26 +0550 2023	N/A	RUNNING	UNDEFINED
<a href="#">application_1675999795986_60672</a>	itv006110	pyspark-shell	SPARK		default	0	Wed Jun 14 18:37:04 +0550 2023	Wed Jun 14 18:37:04 +0550 2023	N/A	RUNNING	UNDEFINED
<a href="#">application_1675999795986_60671</a>	itv007207	pyspark-shell	SPARK		default	0	Wed Jun 14 18:34:17 +0550 2023	Wed Jun 14 18:34:18 +0550 2023	N/A	RUNNING	UNDEFINED
<a href="#">application_1675999795986_60670</a>	itv006826	pyspark-shell	SPARK		default	0	Wed Jun 14 18:32:36 +0550 2023	Wed Jun 14 18:32:36 +0550 2023	N/A	RUNNING	UNDEFINED
<a href="#">application_1675999795986_60669</a>	itv006911	SparkSQL::172.16.1.102	SPARK		default	0	Wed Jun 14 18:28:41 +0550 2023	Wed Jun 14 18:28:41 +0550 2023	N/A	RUNNING	UNDEFINED
<a href="#">application_1675999795986_60668</a>	itv006926	pyspark-shell	SPARK		default	0	Wed Jun 14 18:26:46	Wed Jun 14 18:26:46 +0550	N/A	RUNNING	UNDEFINED

4. In the Application there is FINISHED tab which shows only FINISHED applications.



## FINISHED Applications

Cluster

About  
Nodes  
Node Labels  
Applications  
NEW  
NEW SAVING  
SUBMITTED  
ACCEPTED  
RUNNING  
FINISHED  
FAILED  
KILLED  
Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory T
60106	0	17	60089	51	85 GB	151.00 GB

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes
3	0	0	1


Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>

Show 20 entries

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	Final Status	Run Cor
<a href="#">application_1675999795986_60664</a>	itv007293	pyspark-shell	SPARK		default	0	Wed Jun 14 18:24:46 +0550 2023	Wed Jun 14 18:24:47 +0550 2023	Wed Jun 14 18:25:16 +0550 2023	FINISHED	SUCCEEDED	N/A
<a href="#">application_1675999795986_60663</a>	itv007293	pyspark-shell	SPARK		default	0	Wed Jun 14 18:24:00 +0550 2023	Wed Jun 14 18:24:01 +0550 2023	Wed Jun 14 18:24:42 +0550 2023	FINISHED	SUCCEEDED	N/A
<a href="#">application_1675999795986_60662</a>	itv006054	pyspark-shell	SPARK		default	0	Wed Jun 14 18:23:59 +0550 2023	Wed Jun 14 18:23:59 +0550 2023	Wed Jun 14 18:26:58 +0550 2023	FINISHED	SUCCEEDED	N/A
<a href="#">application_1675999795986_60660</a>	ana007574	pyspark-shell	SPARK		default	0	Wed Jun 14 18:23:22 +0550 2023	Wed Jun 14 18:23:22 +0550 2023	Wed Jun 14 18:26:35 +0550 2023	FINISHED	SUCCEEDED	N/A
<a href="#">application_1675999795986_60659</a>	itv006054	pyspark-shell	SPARK		default	0	Wed Jun 14 18:21:56 +0550 2023	Wed Jun 14 18:21:56 +0550 2023	Wed Jun 14 18:23:51 +0550 2023	FINISHED	SUCCEEDED	N/A
<a href="#">application_1675999795986_60657</a>	itv006926	pyspark-shell	SPARK		default	0	Wed Jun 14 18:15:02 +0550 2023	Wed Jun 14 18:15:02 +0550 2023	Wed Jun 14 18:15:41 +0550 2023	FINISHED	SUCCEEDED	N/A
<a href="#">application_1675999795986_60656</a>	itv002282	INSERT INTO	MAPREDUCE		default	0	Wed Jun 14 18:14:51	Wed Jun 14 18:14:52 +0550	Wed Jun 14 18:15:04	FINISHED	SUCCEEDED	N/A

5. In the Application there is KILLED tab which shows only KILLED applications.



## KILLED Applications

Cluster

About

Nodes

Node Labels

Applications

NEW

SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	60107	Apps Pending	0	Apps Running	17	Apps Completed	60090	Containers Running	52	Memory Used	87 GB	Memory To	151.00 GB
----------------	-------	--------------	---	--------------	----	----------------	-------	--------------------	----	-------------	-------	-----------	-----------

Cluster Nodes Metrics

Active Nodes	3	Decommissioning Nodes	0	Decommissioned Nodes	1	Lost Nodes	0
--------------	---	-----------------------	---	----------------------	---	------------	---

Scheduler Metrics

Scheduler Type	Capacity Scheduler	Scheduling Resource Type	(memory-mb (unit=Mi), vcores)	Minimum Allocation	<memory:1024, vCores: 1>
----------------	--------------------	--------------------------	-------------------------------	--------------------	--------------------------

Show 20 entries

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	Final Status	Runn Contain
application_1675999795986_60648	itv007171	pyspark-shell	SPARK		default	0	Wed Jun 14 17:52:41 +0550 2023	Wed Jun 14 17:52:41 +0550 2023	Wed Jun 14 18:30:10 +0550 2023	KILLED	KILLED	N/A
application_1675999795986_60647	itv006164	Practice	SPARK		default	0	Wed Jun 14 17:51:17 +0550 2023	Wed Jun 14 17:51:18 +0550 2023	Wed Jun 14 18:30:12 +0550 2023	KILLED	KILLED	N/A
application_1675999795986_60645	itv006996	pyspark-shell	SPARK		default	0	Wed Jun 14 17:45:07 +0550 2023	Wed Jun 14 17:45:07 +0550 2023	Wed Jun 14 18:30:13 +0550 2023	KILLED	KILLED	N/A
application_1675999795986_60644	itv006054	pyspark-shell	SPARK		default	0	Wed Jun 14 17:44:08 +0550 2023	Wed Jun 14 17:44:08 +0550 2023	Wed Jun 14 18:15:23 +0550 2023	KILLED	KILLED	N/A
application_1675999795986_60642	itv006082	pyspark-shell	SPARK		default	0	Wed Jun 14 17:41:38 +0550 2023	Wed Jun 14 17:41:38 +0550 2023	Wed Jun 14 18:15:24 +0550 2023	KILLED	KILLED	N/A
application_1675999795986_60637	itv006826	pyspark-shell	SPARK		default	0	Wed Jun 14 17:34:21 +0550 2023	Wed Jun 14 17:34:21 +0550 2023	Wed Jun 14 18:15:24 +0550 2023	KILLED	KILLED	N/A
application_1675999795986_60635	itv005875	pyspark-shell	SPARK		default	0	Wed Jun 14 17:30:35 +0550 2023	Wed Jun 14 17:30:36 +0550 2023	Wed Jun 14 18:15:25 +0550 2023	KILLED	KILLED	N/A

6. In Tools there is Configuration tab which shows all the configuration of the cluster.

Tools

Configuration


Local logs

Server

This XML file does not appear to have any style information associated with it. The document tree is shown below.

```
<configuration>
  <property>
    <name>mapreduce.jobhistory.jhist.format</name>
    <value>binary</value>
    <final>>false</final>
    <source>mapred-default.xml</source>
  </property>
  <property>
    <name>fs.s3a.retry.interval</name>
    <value>500ms</value>
    <final>>false</final>
    <source>core-default.xml</source>
  </property>
  <property>
    <name>hadoop.proxyuser.hive.groups</name>
    <value>*</value>
    <final>>false</final>
    <source>core-site.xml</source>
  </property>
  <property>
    <name>dfs.block.access.token.lifetime</name>
    <value>600</value>
    <final>>false</final>
    <source>hdfs-default.xml</source>
  </property>
  <property>
    <name>mapreduce.job.heap.memory-mb.ratio</name>
    <value>0.8</value>
    <final>>false</final>
    <source>mapred-default.xml</source>
  </property>
  <property>
    <name>mapreduce.map.log.level</name>
    <value>INFO</value>
    <final>>false</final>
    <source>mapred-default.xml</source>
  </property>
  <property>
    <name>dfs.namenode.lazypersist.file.scrub.interval.sec</name>
    <value>300</value>
  </property>
</configuration>
```

7. In Tools there is Local logs tab which shows all the logs of the cluster.



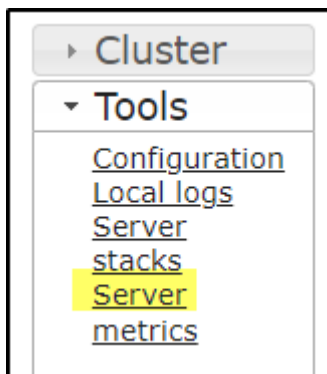
## Directory: /logs/

Name ↕	Last Modified	Size
<a href="#">hadoop-hdfs-resourcemanager-m02.log</a>	Jun 14, 2023 9:14:42 AM	176,530,783 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.1</a>	Jun 7, 2023 7:37:30 AM	268,435,569 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.10</a>	Sep 29, 2022 11:45:10 AM	268,435,614 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.11</a>	Sep 6, 2022 11:15:09 AM	268,435,467 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.12</a>	Aug 2, 2022 2:47:46 PM	268,435,572 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.13</a>	Jul 11, 2022 12:19:33 AM	268,435,485 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.14</a>	Jun 20, 2022 7:14:33 AM	268,435,570 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.15</a>	May 26, 2022 9:38:10 PM	268,435,602 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.16</a>	Apr 27, 2022 12:28:17 PM	268,435,463 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.17</a>	Apr 4, 2022 7:50:27 AM	268,435,481 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.18</a>	Mar 1, 2022 10:21:27 AM	268,435,597 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.19</a>	Jan 28, 2022 8:01:17 AM	268,435,870 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.2</a>	May 24, 2023 7:13:19 AM	268,435,596 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.20</a>	Dec 29, 2021 4:20:51 PM	268,435,747 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.3</a>	May 4, 2023 11:38:36 PM	268,435,778 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.4</a>	Apr 4, 2023 4:16:25 AM	268,435,542 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.5</a>	Mar 1, 2023 12:39:29 AM	268,435,751 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.6</a>	Feb 8, 2023 11:57:42 PM	268,435,817 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.7</a>	Jan 21, 2023 12:44:14 AM	268,435,556 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.8</a>	Dec 17, 2022 8:16:37 AM	268,435,508 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.log.9</a>	Nov 2, 2022 12:08:24 PM	268,435,489 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.out</a>	Feb 9, 2023 10:29:55 PM	2,226 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.out.1</a>	Feb 3, 2023 10:50:27 AM	2,226 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.out.2</a>	Feb 3, 2023 2:45:10 AM	2,219 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.out.3</a>	Feb 2, 2023 10:53:06 PM	2,226 bytes
<a href="#">hadoop-hdfs-resourcemanager-m02.out.4</a>	Feb 2, 2023 7:21:42 PM	2,219 bytes

```
2023-06-07 07:37:30,799 INFO org.apache.hadoop.yarn.server.resourcemanager.rmapp.RMAppImpl: application_1675999795986_55865 State change from SUBMITTED to ACCEPTED on event = APP_ACCEPTED
2023-06-07 07:37:30,799 INFO org.apache.hadoop.yarn.server.resourcemanager.ApplicationMasterService: Registering app attempt : appattempt_1675999795986_55865_000001
2023-06-07 07:37:30,799 INFO org.apache.hadoop.yarn.server.resourcemanager.rmapp.attempt.RMAppAttemptImpl: appattempt_1675999795986_55865_000001 State change from NEW to SUBMITTED on event = START
2023-06-07 07:37:30,800 INFO org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.LeafQueue: Application application_1675999795986_55865 from user: itv006115 activated in queue: root.default
2023-06-07 07:37:30,800 INFO org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.LeafQueue: Application added - appId: application_1675999795986_55865 user: itv006115, leaf-queue: root.default #user-pending-applications: 0 #user-active-applications: 1 #queue-pending-applications: 0 #queue-active-applications: 14
2023-06-07 07:37:30,800 INFO org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.CapacityScheduler: Added Application Attempt appattempt_1675999795986_55865_000001 to scheduler from user itv006115 in queue root.default
2023-06-07 07:37:30,800 INFO org.apache.hadoop.yarn.server.resourcemanager.rmapp.attempt.RMAppAttemptImpl: appattempt_1675999795986_55865_000001 State change from SUBMITTED to SCHEDULED on event = ATTEMPT_ADDED
2023-06-07 07:37:31,020 INFO org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.allocation.AbstractContainerAllocator: assignedContainer application attempt=appattempt_1675999795986_55865_000001 container=null
queue=default ClusterResource<memory:154620, vCores:90> type=OFF_SWITCH requestedPartition=
2023-06-07 07:37:31,020 INFO org.apache.hadoop.yarn.server.resourcemanager.rmcontainer.RMContainerImpl: container_1675999795986_55865_01_000001 Container Transitioned from NEW to ALLOCATED
2023-06-07 07:37:31,020 INFO org.apache.hadoop.yarn.server.resourcemanager.scheduler.common.fica.FicaSchedulerNode: Assigned container container_1675999795986_55865_01_000001 of capacity <memory:1024, vCores:1> on host w03.itversity.com:41791, which has 14 containers, <memory:24576, vCores:14> used and <memory:26964, vCores:16> available after allocation
2023-06-07 07:37:31,020 INFO org.apache.hadoop.yarn.server.resourcemanager.security.NMTokenSecretManagerInRM: Sending NMToken for nodeid : w03.itversity.com:41791 for container : container_1675999795986_55865_01_000001
2023-06-07 07:37:31,020 INFO org.apache.hadoop.yarn.server.resourcemanager.rmcontainer.RMContainerImpl: container_1675999795986_55865_01_000001 Container Transitioned from ALLOCATED to ACQUIRED
2023-06-07 07:37:31,021 INFO org.apache.hadoop.yarn.server.resourcemanager.security.NMTokenSecretManagerInRM: Clear node set for appattempt_1675999795986_55865_000001
2023-06-07 07:37:31,021 INFO org.apache.hadoop.yarn.server.resourcemanager.rmapp.attempt.RMAppAttemptImpl: Storing attempt: AppId: application_1675999795986_55865 AttemptId: appattempt_1675999795986_55865_000001 MasterContainer: Container: [ContainerId: container_1675999795986_55865_01_000001, AllocationRequestId: -1, Version: 0, NodeId: w03.itversity.com:41791, NodeHttpAddress: w03.itversity.com:8042, <memory:1024, vCores:1>, Priority: 0, Token: Token { kind: ContainerToken, service: 172.16.1.107:41791 }, ExecutionType: GUARANTEED, ]
2023-06-07 07:37:31,021 INFO org.apache.hadoop.yarn.server.resourcemanager.rmapp.attempt.RMAppAttemptImpl: appattempt_1675999795986_55865_000001 State change from SCHEDULED to ALLOCATED_SAVING on event = CONTAINER_ALLOCATED
2023-06-07 07:37:31,021 INFO org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.ParentQueue: assignedContainer queue=root usedCapacity=0.6622688 absoluteUsedCapacity=0.6622688 used=<memory:102400, vCores:56> cluster=<memory:154620, vCores:90>
2023-06-07 07:37:31,021 INFO org.apache.hadoop.yarn.server.resourcemanager.rmapp.attempt.RMAppAttemptImpl: appattempt_1675999795986_55865_000001 State change from ALLOCATED_SAVING to ALLOCATED on event = ATTEMPT_NEW_SAVED
2023-06-07 07:37:31,021 INFO org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.CapacityScheduler: Allocation proposal accepted
2023-06-07 07:37:31,021 INFO org.apache.hadoop.yarn.server.resourcemanager.amlauncher.AMLauncher: Launching masterappattempt_1675999795986_55865_000001
2023-06-07 07:37:31,022 INFO org.apache.hadoop.yarn.server.resourcemanager.amlauncher.AMLauncher: Setting up container Container: [ContainerId: container_1675999795986_55865_01_000001, AllocationRequestId: -1, Version: 0, NodeId: w03.itversity.com:41791, NodeHttpAddress: w03.itversity.com:8042, Resource: <memory:1024, vCores:1>, Priority: 0, Token: Token { kind: ContainerToken, service: 172.16.1.107:41791 }, ExecutionType: GUARANTEED, ] for AM appattempt_1675999795986_55865_000001
2023-06-07 07:37:31,022 INFO org.apache.hadoop.yarn.server.resourcemanager.security.AMRMTokensSecretManager: Create AMRMToken for ApplicationAttempt: appattempt_1675999795986_55865_000001
2023-06-07 07:37:31,022 INFO org.apache.hadoop.yarn.server.resourcemanager.security.AMRMTokensSecretManager: Creating password for appattempt_1675999795986_55865_000001
2023-06-07 07:37:31,027 INFO org.apache.hadoop.yarn.server.resourcemanager.amlauncher.AMLauncher: Done launching container Container: [ContainerId: container_1675999795986_55865_01_000001, AllocationRequestId: -1, Version: 0, NodeId: w03.itversity.com:41791, NodeHttpAddress: w03.itversity.com:8042, Resource: <memory:1024, vCores:1>, Priority: 0, Token: Token { kind: ContainerToken, service: 172.16.1.107:41791 }, ExecutionType: GUARANTEED, ] for AM appattempt_1675999795986_55865_000001
2023-06-07 07:37:31,027 INFO org.apache.hadoop.yarn.server.resourcemanager.rmapp.attempt.RMAppAttemptImpl: appattempt_1675999795986_55865_000001 State change from ALLOCATED to LAUNCHED on event = LAUNCHED
2023-06-07 07:37:31,027 INFO org.apache.hadoop.yarn.server.resourcemanager.rmapp.RMAppImpl: update the launch time for applicationId: application_1675999795986_55865, attemptId: appattempt_1675999795986_55865_000001 launchTime: 1686137851027
2023-06-07 07:37:31,027 INFO org.apache.hadoop.yarn.server.resourcemanager.recovery.RMStateStore: Updating info for app: application_1675999795986_55865
2023-06-07 07:37:32,021 INFO org.apache.hadoop.yarn.server.resourcemanager.rmcontainer.RMContainerImpl: container_1675999795986_55865_01_000001 Container Transitioned from ACQUIRED to RUNNING
2023-06-07 07:37:32,793 INFO org.apache.hadoop.yarn.server.resourcemanager.DefaultAMSPProcessor: AM registration appattempt_1675999795986_55865_000001 (auth:SIMPLE)
2023-06-07 07:37:32,798 INFO org.apache.hadoop.yarn.server.resourcemanager.DefaultAMSPProcessor: AM registration appattempt_1675999795986_55865_000001
2023-06-07 07:37:32,888 INFO org.apache.hadoop.yarn.server.resourcemanager.rmapp.attempt.RMAppAttemptImpl: appattempt_1675999795986_55865_000001 State change from LAUNCHED to RUNNING on event = REGISTERED
2023-06-07 07:37:32,888 INFO org.apache.hadoop.yarn.server.resourcemanager.rmapp.RMAppImpl: application_1675999795986_55865 State change from ACCEPTED to RUNNING on event = ATTEMPT_REGISTERED
```



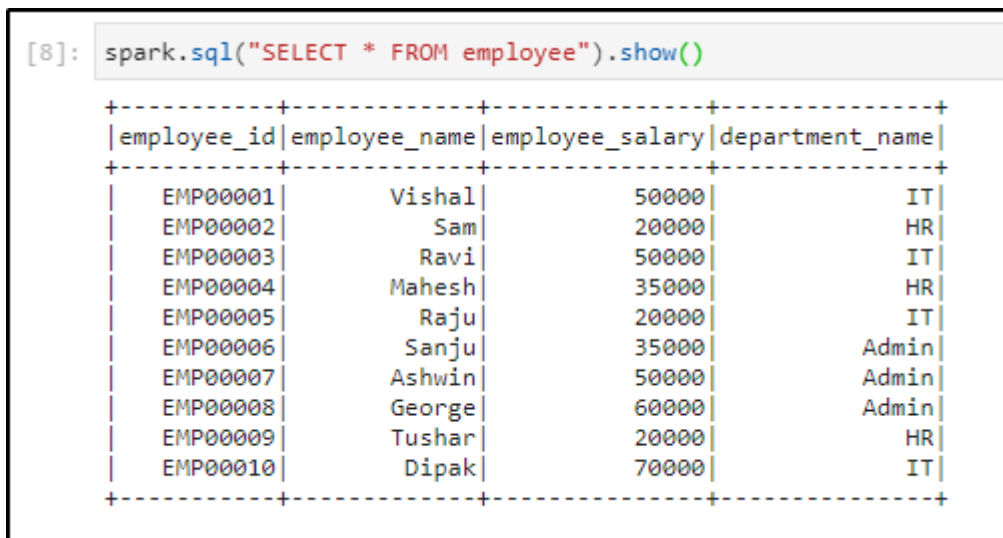
8. In Tools there is Server tab which shows server details of the cluster in json format.



### Question 3.

Perform various queries

Ans: - I have choose below employee dataset



### 1. Running salary of employee

#### Running total Salary

```
[36]: spark.sql("""SELECT employee_id,employee_name,employee_salary,department_name
,SUM(employee_salary) OVER(ORDER BY employee_id) AS running_salary
FROM employee""").show()
```

employee_id	employee_name	employee_salary	department_name	running_salary
EMP00001	Vishal	50000	IT	50000
EMP00002	Sam	20000	HR	70000
EMP00003	Ravi	50000	IT	120000
EMP00004	Mahesh	35000	HR	155000
EMP00005	Raju	20000	IT	175000
EMP00006	Sanju	35000	Admin	210000
EMP00007	Ashwin	50000	Admin	260000
EMP00008	George	60000	Admin	320000
EMP00009	Tushar	20000	HR	340000
EMP00010	Dipak	70000	IT	410000

### 2. Total salary of each department

#### Total Salary of each Department

```
[37]: spark.sql("""SELECT department_name,SUM(employee_salary) AS total_salary
FROM employee GROUP BY department_name""").show()
```

department_name	total_salary
HR	75000
Admin	145000
IT	190000

### 3. Second Highest salary of each department

#### Second Highest salary of each department

```
[12]: spark.sql("""WITH CTESecondSalary AS (
SELECT employee_id,employee_name,employee_salary,department_name
,dense_rank(employee_salary) OVER(PARTITION BY department_name ORDER BY employee_salary DESC) AS denserank_salary
FROM employee
)
SELECT employee_id,employee_name,employee_salary,department_name
FROM CTESecondSalary
WHERE denserank_salary = 2
""").show()
```

employee_id	employee_name	employee_salary	department_name
EMP00002	Sam	20000	HR
EMP00009	Tushar	20000	HR
EMP00007	Ashwin	50000	Admin
EMP00001	Vishal	50000	IT
EMP00003	Ravi	50000	IT

#### 4. Percentage salary of each employee by department

```
[18]: spark.sql("""WITH CTEmployee AS (
SELECT employee_id,employee_name,employee_salary,department_name
,SUM(employee_salary) OVER(PARTITION BY department_name) AS total_salary
FROM employee
)
SELECT employee_id,employee_name,employee_salary,department_name,ROUND((employee_salary/total_salary)*100,2) AS percent_salary
FROM CTEmployee
""").show()
```

employee_id	employee_name	employee_salary	department_name	percent_salary
EMP00002	Sam	20000	HR	26.67
EMP00004	Mahesh	35000	HR	46.67
EMP00009	Tushar	20000	HR	26.67
EMP00006	Sanju	35000	Admin	24.14
EMP00007	Ashwin	50000	Admin	34.48
EMP00008	George	60000	Admin	41.38
EMP00001	Vishal	50000	IT	26.32
EMP00003	Ravi	50000	IT	26.32
EMP00005	Raju	20000	IT	10.53
EMP00010	Dipak	70000	IT	36.84

#### 5. Difference of Salary between current and next highest salary of employee

Difference of Salary between current and next highest salary of employee

```
[25]: spark.sql("""WITH CTEmployee AS (
SELECT employee_id,employee_name,employee_salary,department_name
,LAG(employee_salary) OVER(PARTITION BY department_name ORDER BY employee_salary) AS previous_salary
FROM employee
)
SELECT employee_id,employee_name,employee_salary,department_name,IFNULL(employee_salary - previous_salary,0) AS salary_diff
FROM CTEmployee
""").show()
```

employee_id	employee_name	employee_salary	department_name	salary_diff
EMP00002	Sam	20000	HR	0
EMP00009	Tushar	20000	HR	0
EMP00004	Mahesh	35000	HR	15000
EMP00006	Sanju	35000	Admin	0
EMP00007	Ashwin	50000	Admin	15000
EMP00008	George	60000	Admin	10000
EMP00005	Raju	20000	IT	0
EMP00001	Vishal	50000	IT	30000
EMP00003	Ravi	50000	IT	0
EMP00010	Dipak	70000	IT	20000

#### 6. Pivot for employee with department

Pivot for employee with department\_name

```
[35]: spark.sql("""SELECT employee_id,employee_name,IFNULL(HR,0) AS HR,IFNULL(Admin,0) AS Admin,IFNULL(IT,0) AS IT
FROM employee
PIVOT (
MIN(employee_salary)
FOR department_name IN ('HR','Admin','IT')
)""").show()
```

employee_id	employee_name	HR	Admin	IT
EMP00007	Ashwin	0	50000	0
EMP00006	Sanju	0	35000	0
EMP00010	Dipak	0	0	70000
EMP00004	Mahesh	35000	0	0
EMP00008	George	0	60000	0
EMP00005	Raju	0	0	20000
EMP00009	Tushar	20000	0	0
EMP00001	Vishal	0	0	50000
EMP00002	Sam	20000	0	0
EMP00003	Ravi	0	0	50000

## Question 4.

### How to deal with nulls in Apache Spark?

Ans: -

1. In Spark, if there is datatype mismatch then due to default mode permissive. it will display as NULL.

```
[4]: df_orders_03.show()
```

order_id	order_date	cust_id	order_status
1	2013-07-25	11599	CLOSED
2	2013-07-25	256	PENDING_PAYMENT
3	2013-07-25	12111	COMPLETE
4	2013-07-25	8827	CLOSED
5	2013-07-25	11318	COMPLETE
6	2013-07-25	7130	COMPLETE
7	2013-07-25	null	COMPLETE
8	2013-07-25	2911	PROCESSING
9	2013-07-25	null	PENDING_PAYMENT
10	2013-07-25	5648	PENDING_PAYMENT

We can change this behaviour by changing the mode to failfast or dropmalformed.

2. While defining schema we can set the column as not allow NULL then if we try to insert NULL in that column it will fail.
3. In User define function if we encounter NULL and due to that if code will fail then we can filter the NULL.
4. We can replace NULL values with fill() function in dataframe.
5. For integer column we can replace it with 0

Pivot for employee with department\_name

```
[35]: spark.sql("""SELECT employee_id,employee_name,IFNULL(HR,0) AS HR,IFNULL(Admin,0) AS Admin,IFNULL(IT,0) AS IT
FROM employee
PIVOT (
    MIN(employee_salary)
    FOR department_name IN ('HR','Admin','IT')
)""").show()
```

employee_id	employee_name	HR	Admin	IT
EMP00007	Ashwin	0	50000	0
EMP00006	Sanju	0	35000	0
EMP00010	Dipak	0	0	70000
EMP00004	Mahesh	35000	0	0
EMP00008	George	0	60000	0
EMP00005	Raju	0	0	20000
EMP00009	Tushar	20000	0	0
EMP00001	Vishal	0	0	50000
EMP00002	Sam	20000	0	0
EMP00003	Ravi	0	0	50000